國立陽明交通大學
NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Biological Databases: Theories and Practice 430032

Database Schema and Entity-Relationship Model

Instructor: Prof. LEE, Tzong-Yi (李宗夷)
Email: leetzongyi@nycu.edu.ctw

Professor
Institute of Bioinformatics and Systems Biology,
National Yang Ming Chiao Tung University

Outline

- Modeling
- Constraints
- E-R Diagram
- Weak Entity Sets
- Reduction to Relation Schemas
- Design Issues
- Extended E-R Features

Modeling

- A *database* can be modeled as:
  - a collection of entities,
  - relationship among entities.
- An **entity** is an object that exists and is distinguishable from other objects.
  - Example: specific person, company, event, plant
- Entities have **attributes**
  - Example: people have *names* and *addresses*
- An **entity set** is a set of entities of the same type that share the same properties.
  - Example: set of all persons, companies, trees, holidays

Entity Sets: *instructor* and *student*

instructor_ID	instructor_name
76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor

student-ID	student_name
98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student

Relationship Sets

- A **relationship** is an association among several entities

  Example:

  44553 (Peltier) — *advisor* — 22222 (Einstein)

  student entity — relationship set — *instructor* entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \ldots e_n) \mid e_1 \in E_1, e_2 \in E_2, \ldots, e_n \in E_n\}$$

  where $(e_1, e_2, \ldots, e_n)$ is a relationship

  - Example:

$$(44553,22222) \in advisor$$

Relationship Set: *advisor*

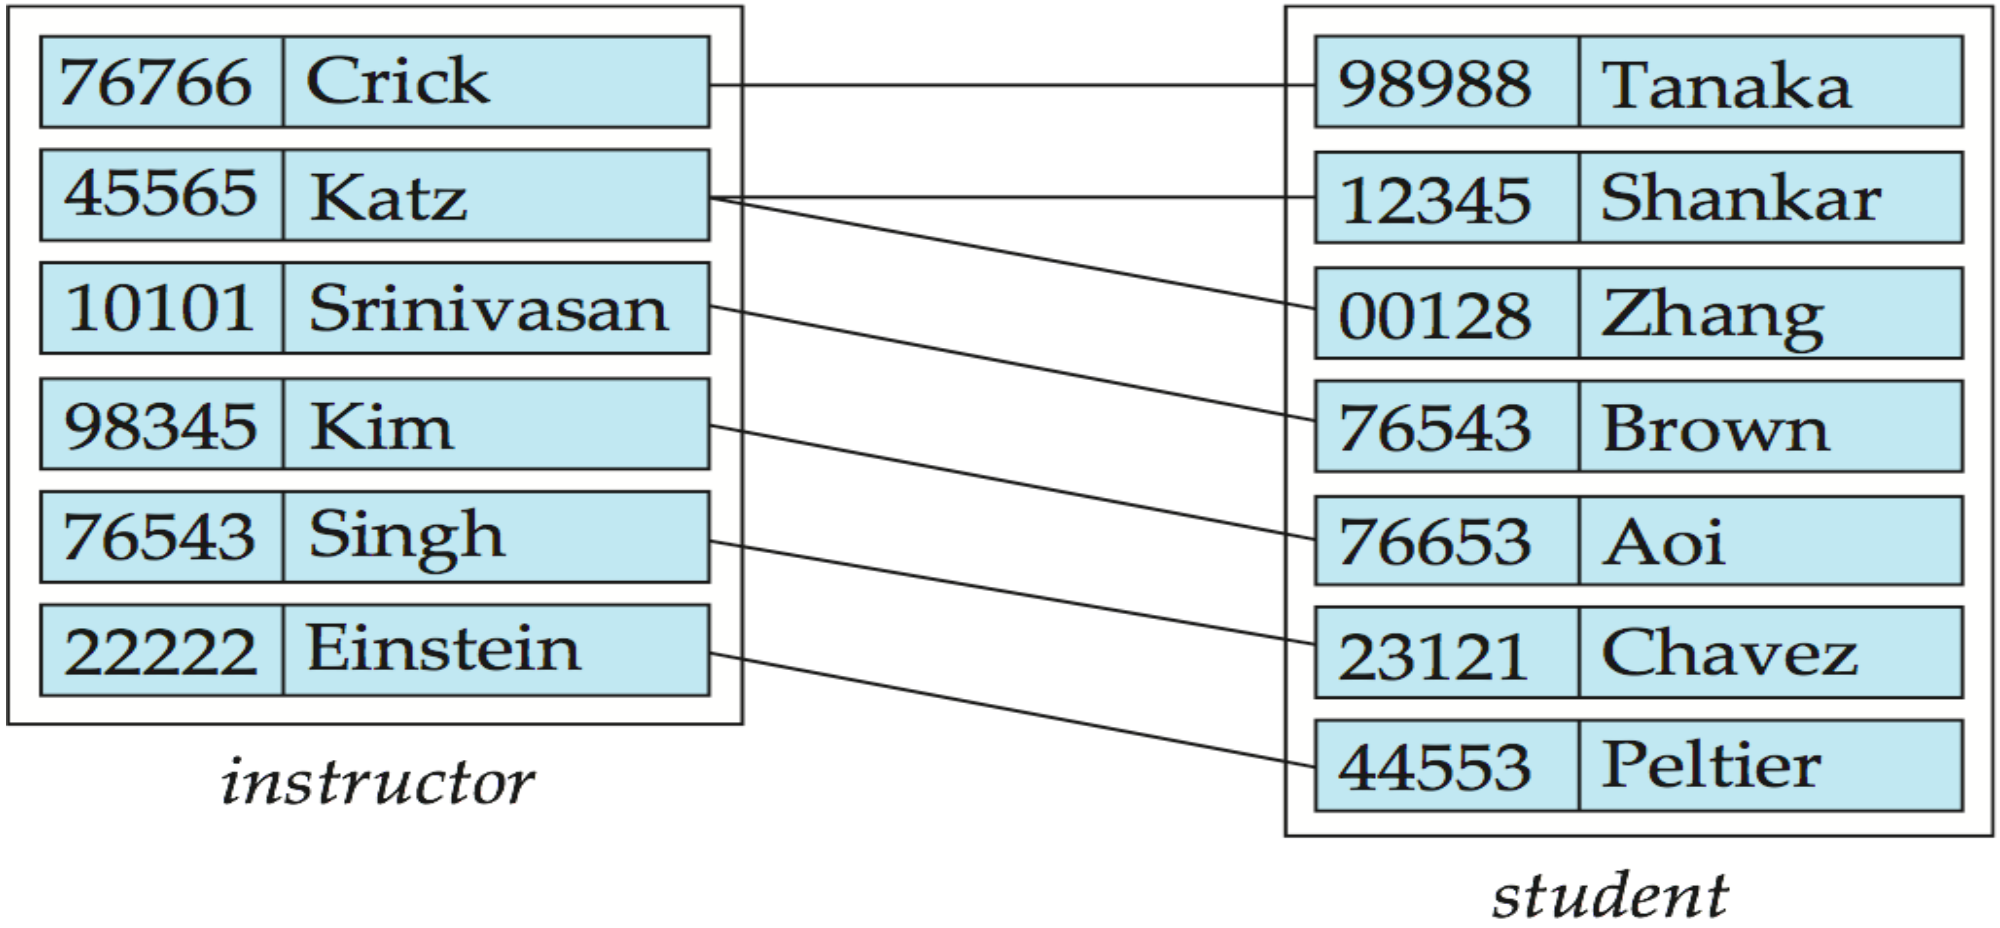

Relationship Sets (Cont.)

- An **attribute** can also be property of a relationship set.
- For instance, the ***advisor*** relationship set between entity sets ***instructor*** and ***student*** may have the attribute ***date*** which tracks when the student started being associated with the advisor

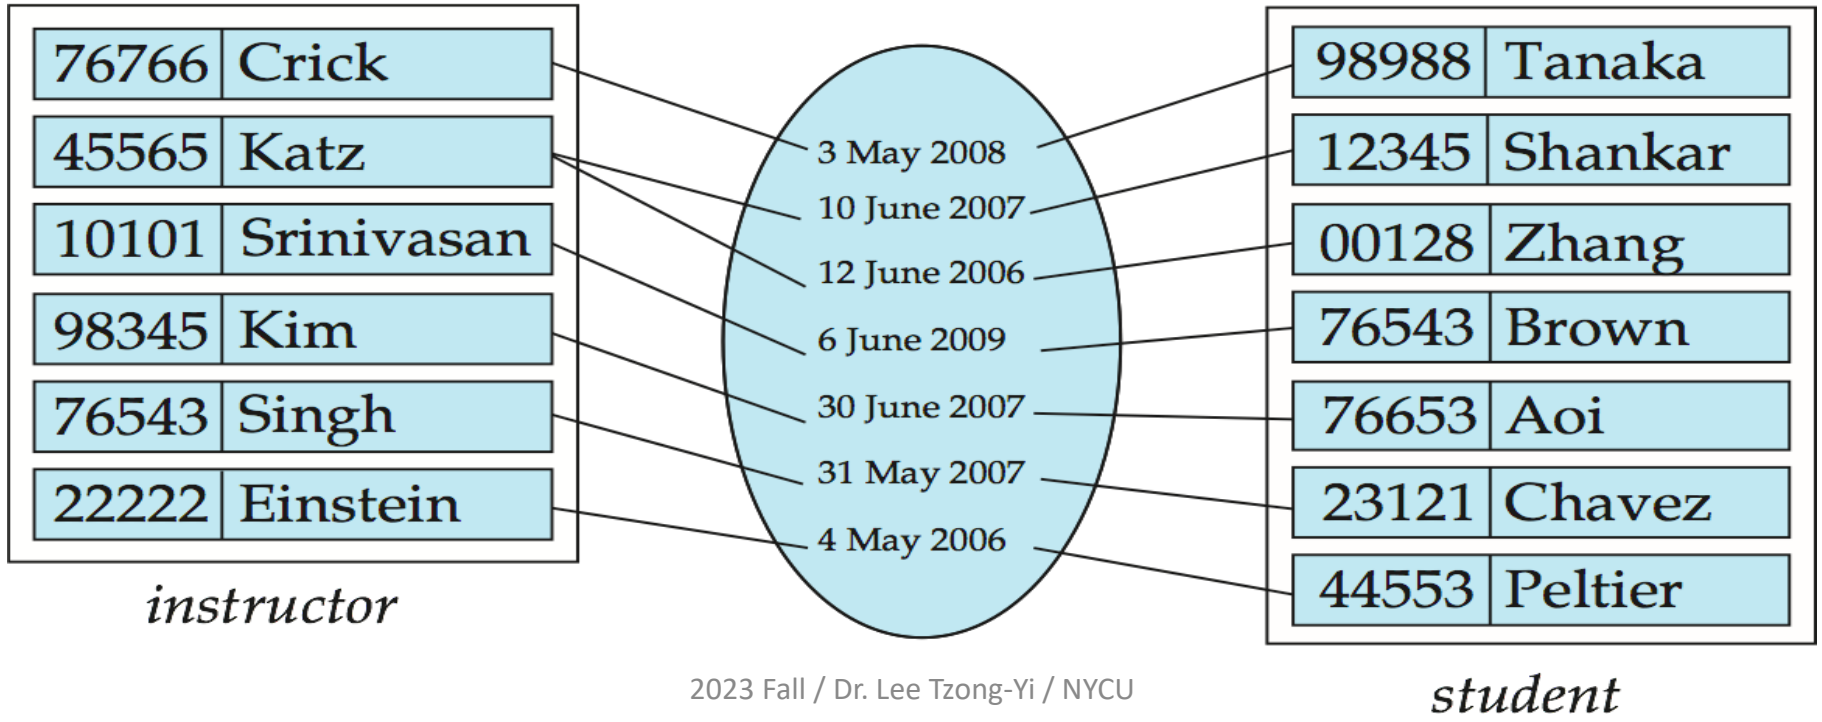

Relationship Sets with Attributes

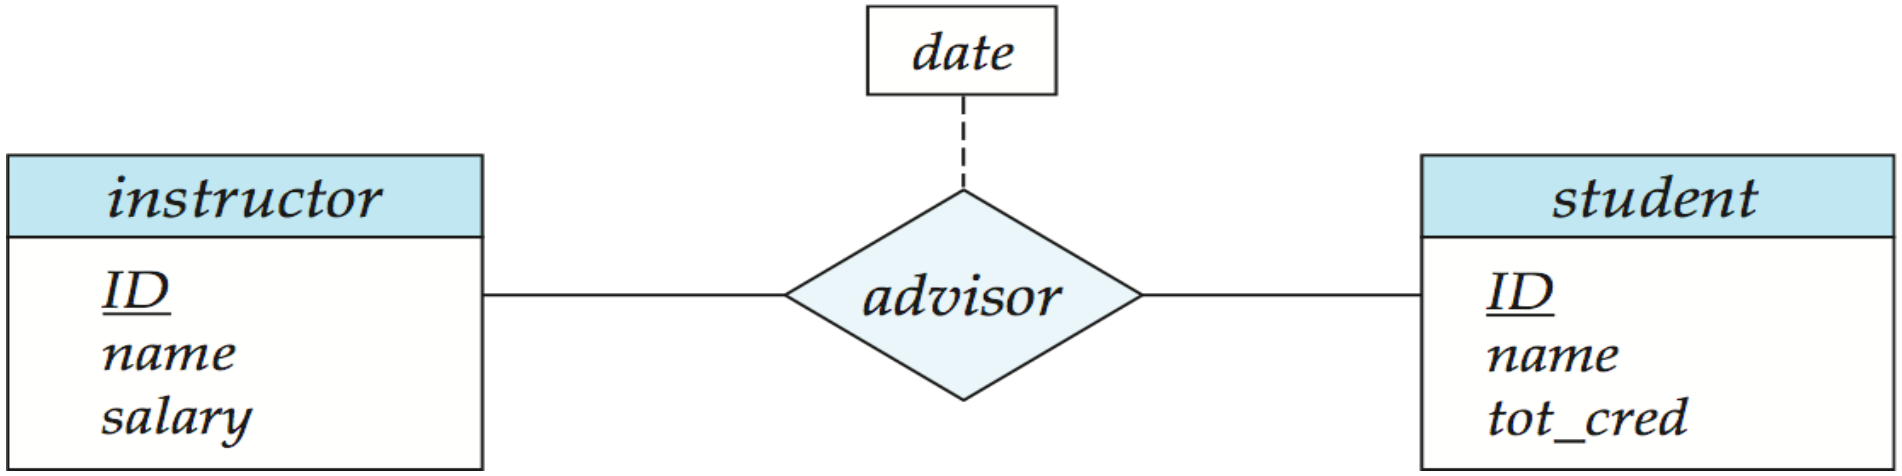

Degree of a Relationship Set

- **Binary relationship**
  - involve two entity sets (or degree two).
  - most relationship sets in a database system are binary.
- Relationships between more than two entity sets are rare. Most relationships are binary.
- Example: *students* work on research *projects* under the guidance of an *instructor*.
  - relationship ***proj_guide*** is a **ternary relationship** between *instructor, student,* and *project*

E-R Diagram with a Ternary Relationship

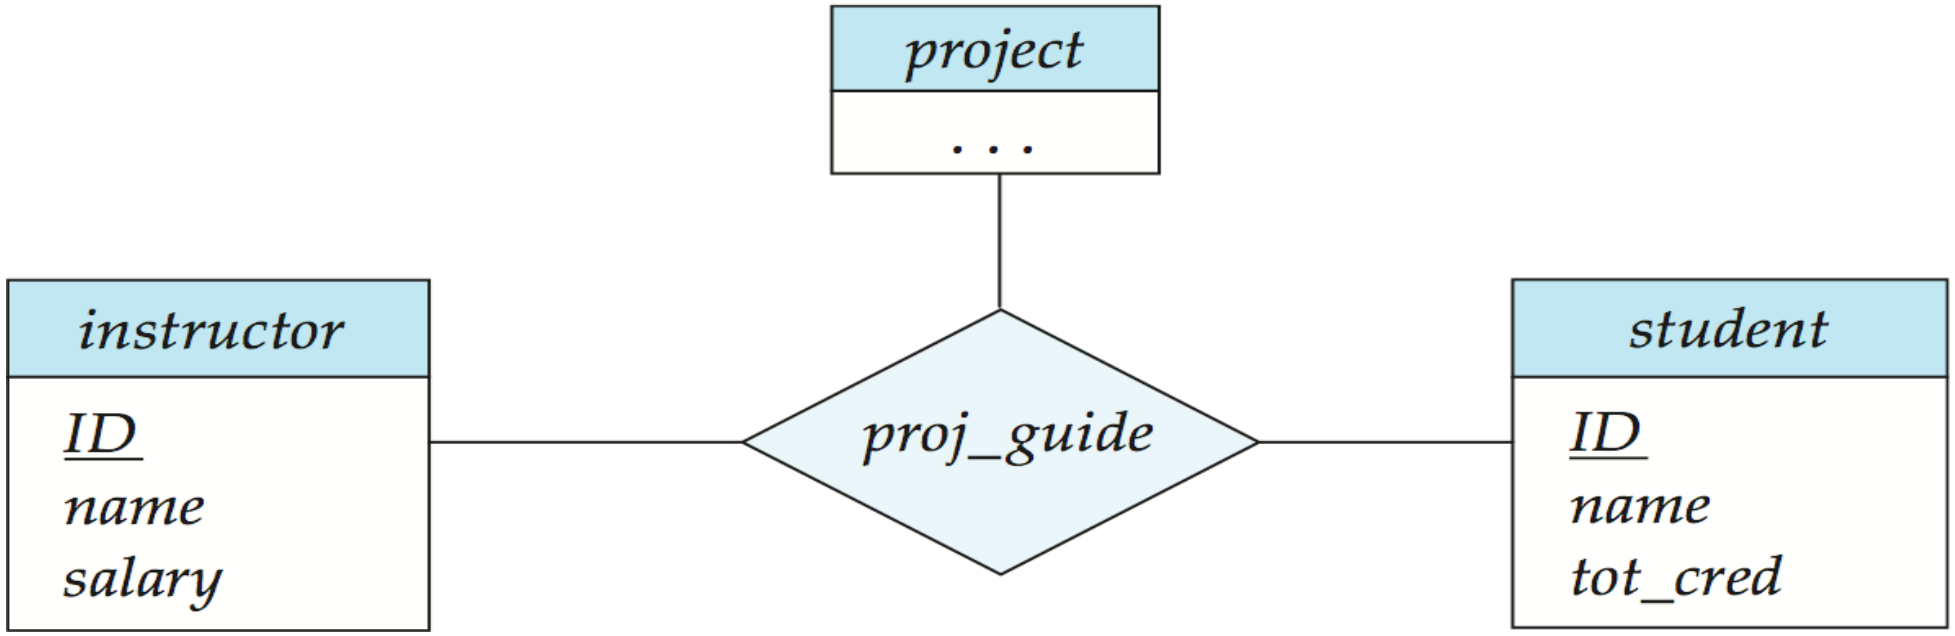

Attributes

- An entity is represented by a set of attributes, that is descriptive properties possessed by all members of an entity set.
  - Example:

    instructor = (*ID, name, street, city, salary*)
    course= (*course_id, title, credits*)
- **Domain** – the set of permitted values for each attribute
- Attribute types:
  - **Simple** and **composite** attributes.
  - **Single-valued** and **multivalued** attributes
    - Example: multivalued attribute: *phone_numbers*
  - **Derived** attributes
    - Can be computed from other attributes
    - Example: ***age*** can be obtained from ***date_of_birth***

Composite Attributes

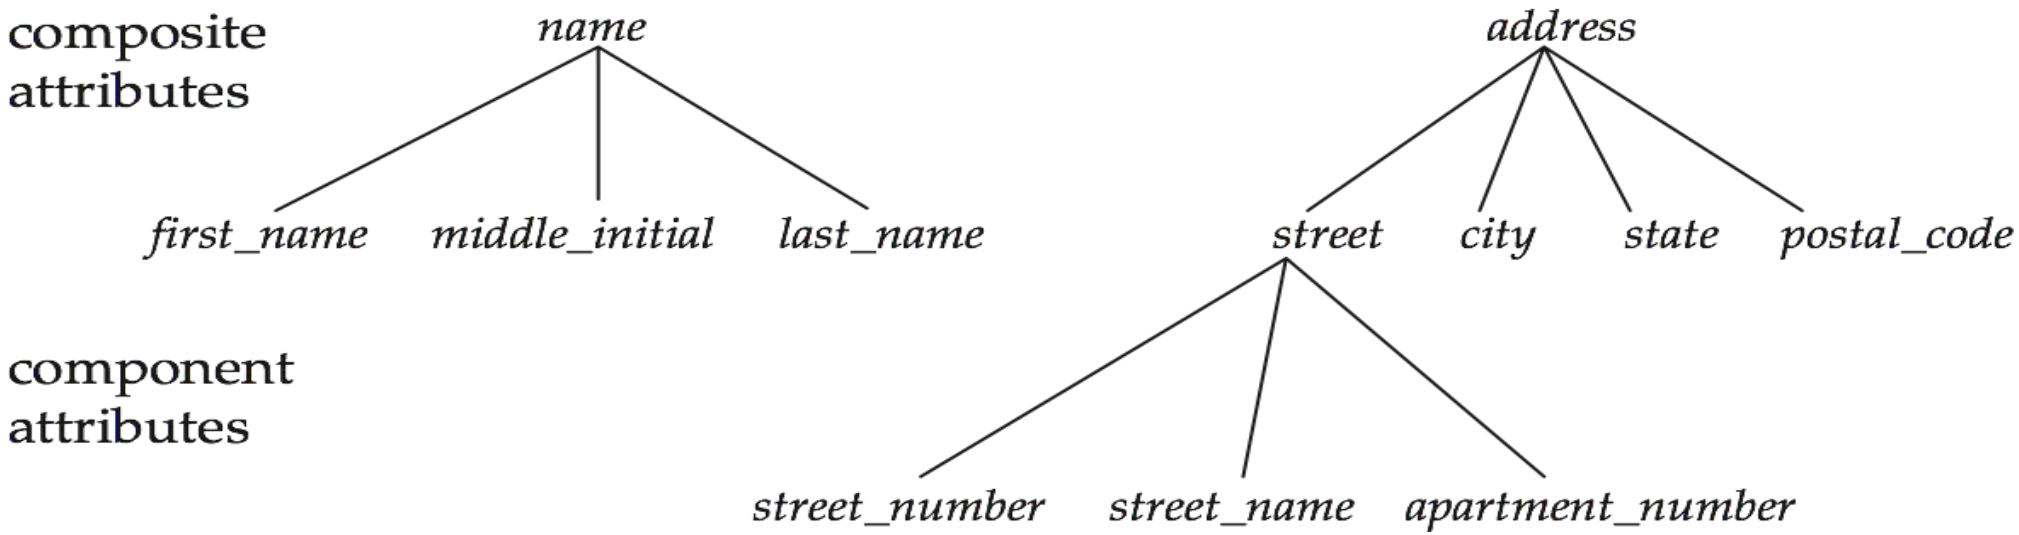

Example

The entity ***instructor*** contains composite, multivalued, and derived attributes.

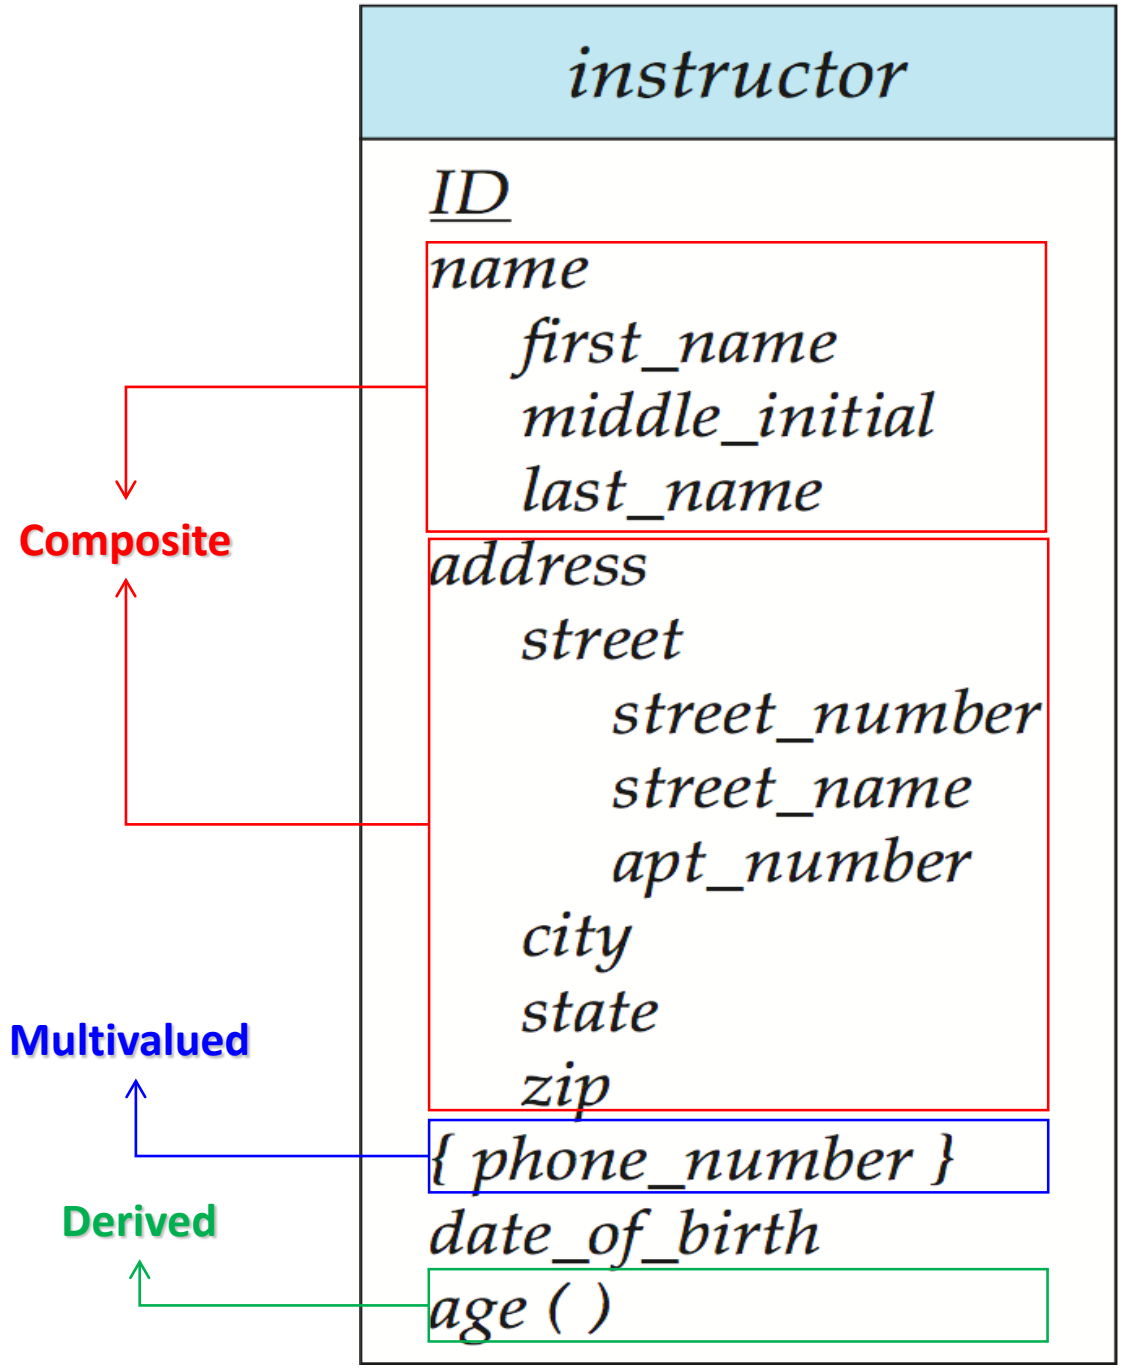

Keys

- A **super key** of an entity set is a set of one or more attributes whose values uniquely determine each entity.
- A **candidate key** of an entity set is a minimal super key
  - *ID* is candidate key of *instructor*
  - *course_id* is candidate key of *course*
- Although several candidate keys may exist, one of the candidate keys is selected to be the **primary key**.

E-R Diagrams

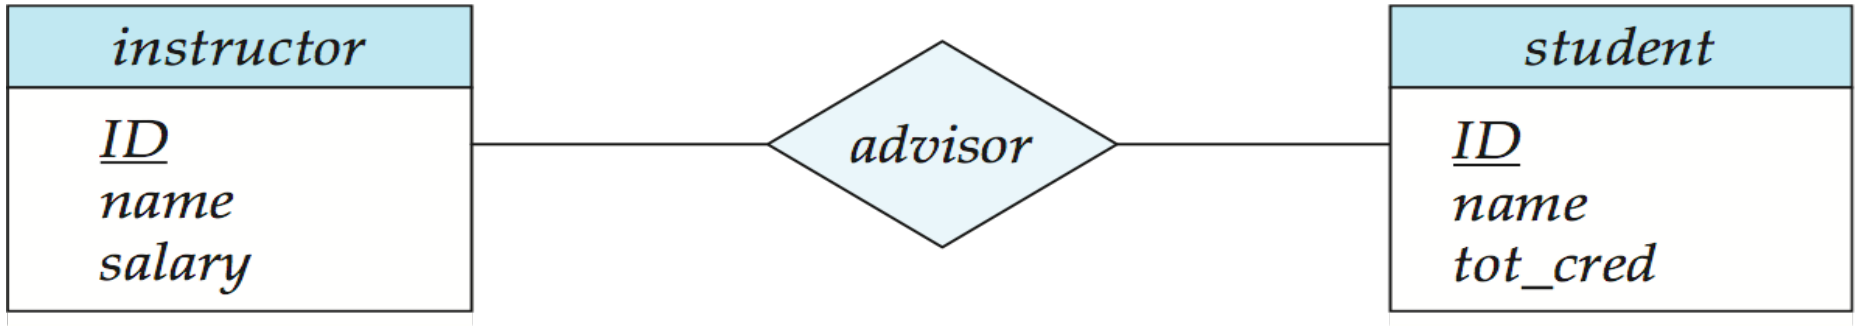

- **Rectangles** represent **entity sets**.
- **Diamonds** represent **relationship sets**.
- **Attributes** listed inside entity rectangle
- **Underline** indicates **primary key** attributes

Roles

- We indicate roles in E-R diagram by labeling the lines that connect diamonds to rectangles.
- Entity sets of a relationship need not be distinct
  - Each occurrence of an entity set plays a "**role**" in the relationship
- The labels "***course_id***" and "***prereq_id***" are called **roles**.

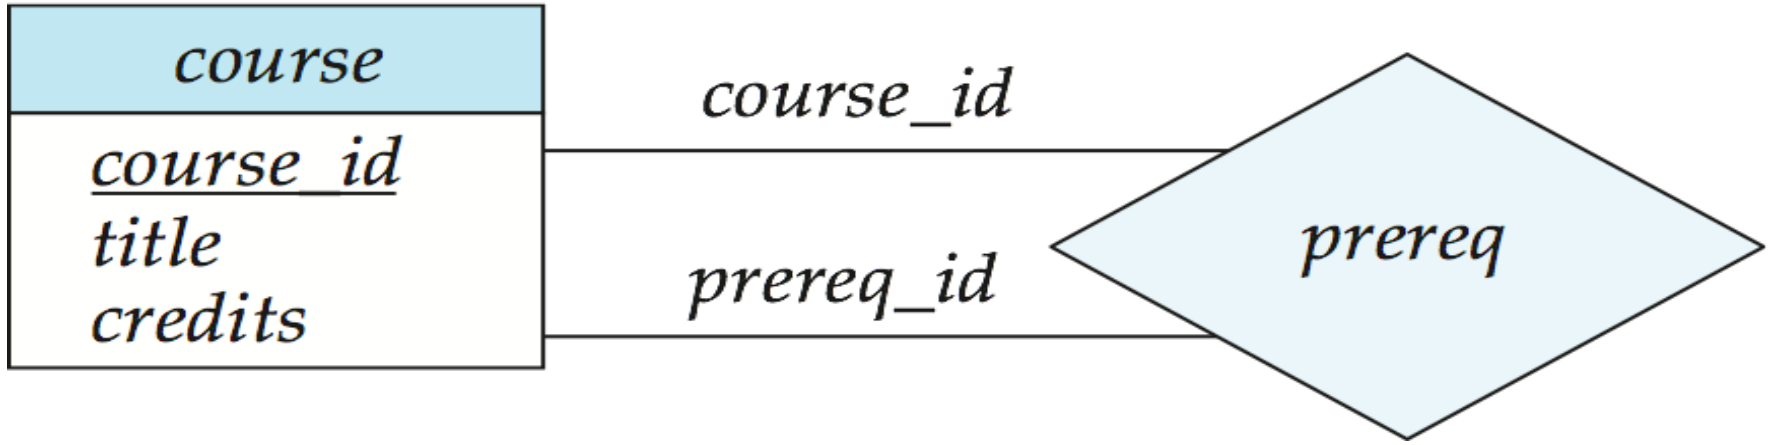

Mapping Cardinality Constraints

- Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
  - **One to one**
  - **One to many**
  - **Many to one**
  - **Many to many**

Mapping Cardinalities

One to one

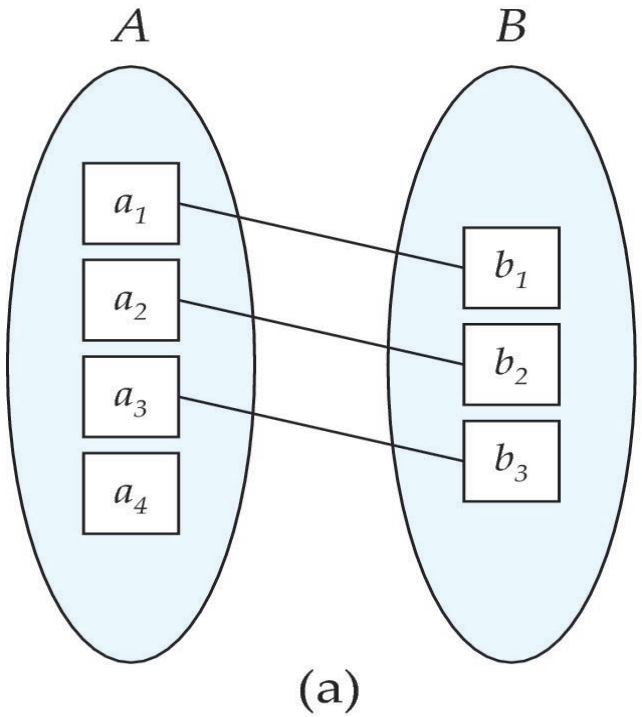

One to many

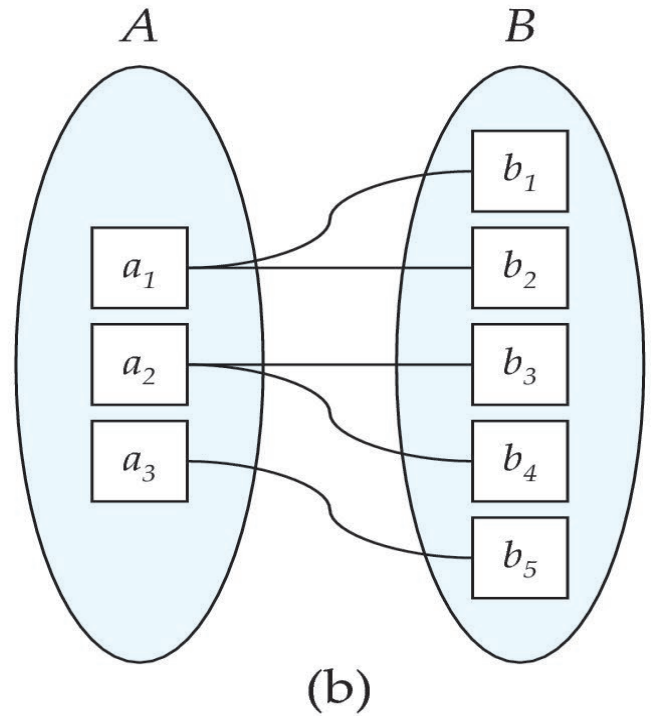

Note: Some elements in *A* and *B* may not be mapped to any elements in the other set

Mapping Cardinalities

Many to one

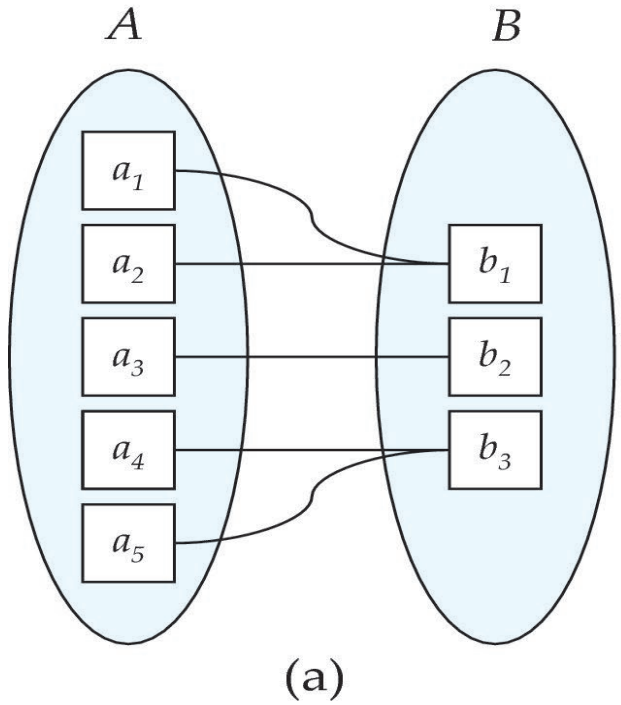

Many to many

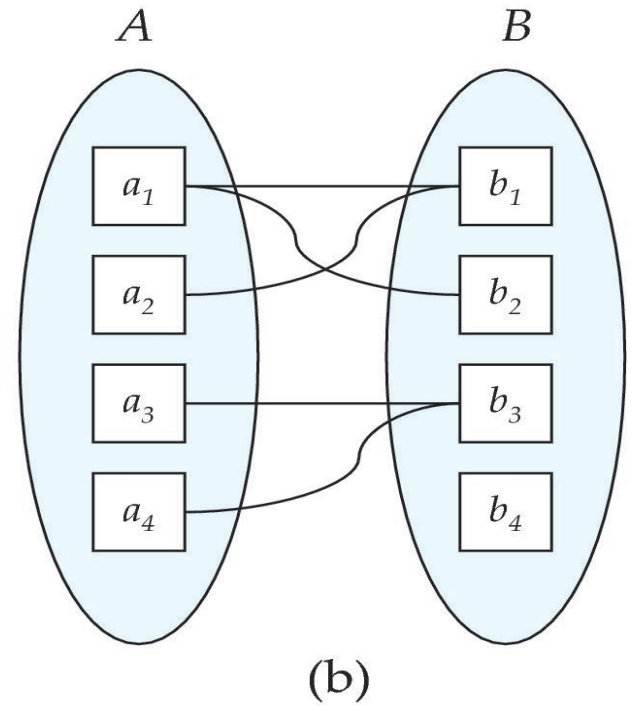

Note: Some elements in *A* and *B* may not be mapped to any elements in the other set

Cardinality Constraints

- We express cardinality constraints by drawing either a directed line (→), signifying "**one**," or an undirected line (—), signifying "**many**," between the relationship set and the entity set.
- One-to-one relationship:
  - A ***student*** is associated with at most one ***instructor*** via the relationship ***advisor***
  - A ***student*** is associated with at most one ***department*** via relationship ***stud_dept***

One-to-One Relationship

- One-to-one relationship between an ***instructor*** and a ***student***
  - an ***instructor*** is associated with at most one ***student*** via ***advisor***
  - and a ***student*** is associated with at most one ***instructor*** via ***advisor***

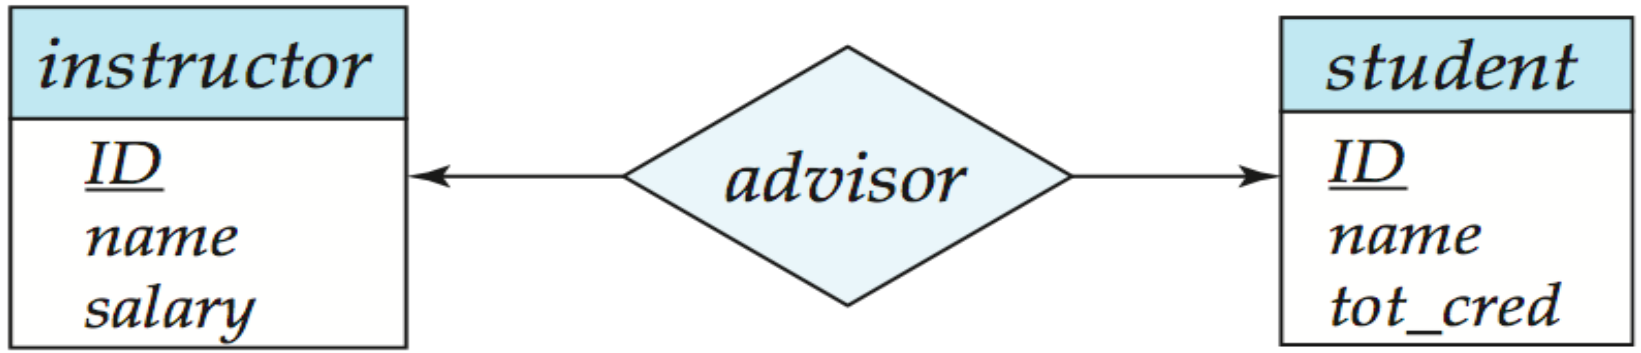

One-to-Many Relationship

- One-to-many relationship between an ***instructor*** and a ***student***
  - an ***instructor*** is associated with several (including 0) ***students*** via ***advisor***
  - a ***student*** is associated with at most one ***instructor*** via ***advisor***

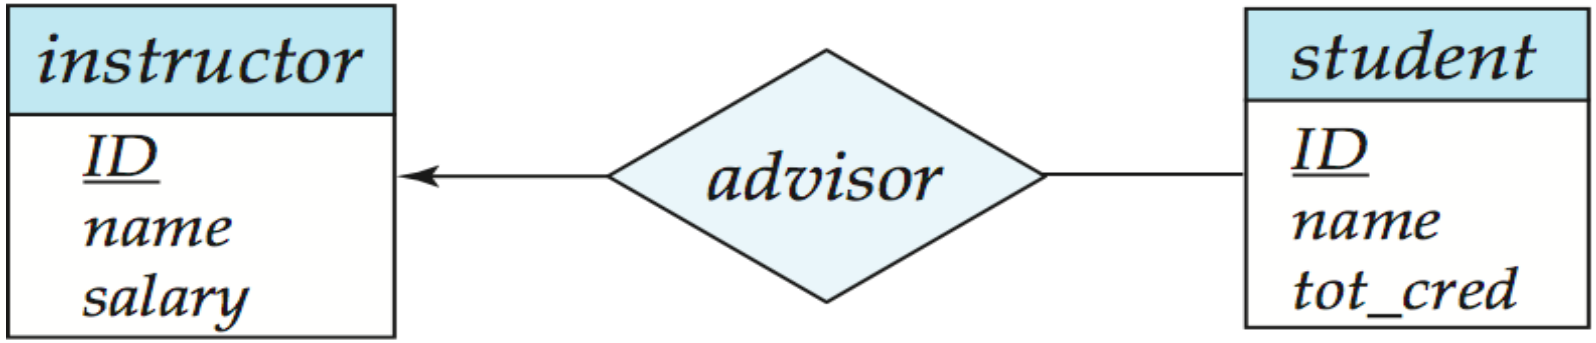

Many-to-One Relationships

- In a many-to-one relationship between an ***instructor*** and a ***student***
  - an ***instructor*** is associated with at most one ***student*** via ***advisor***,
  - and a ***student*** is associated with several (including 0) ***instructors*** via ***advisor***

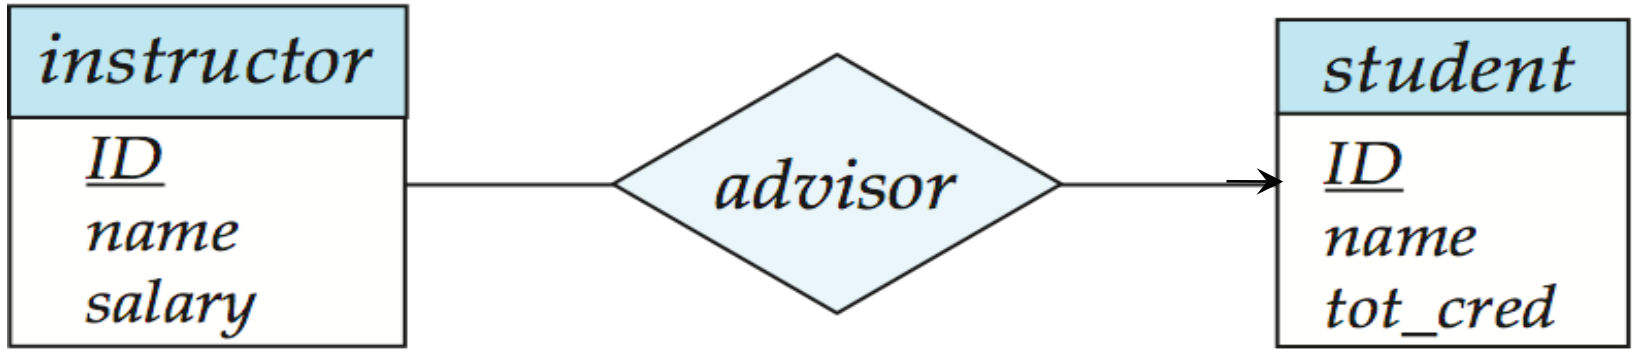

Many-to-Many Relationship

- An ***instructor*** is associated with several (possibly 0) ***students*** via ***advisor***
- A ***student*** is associated with several (possibly 0) ***instructors*** via **advisor**

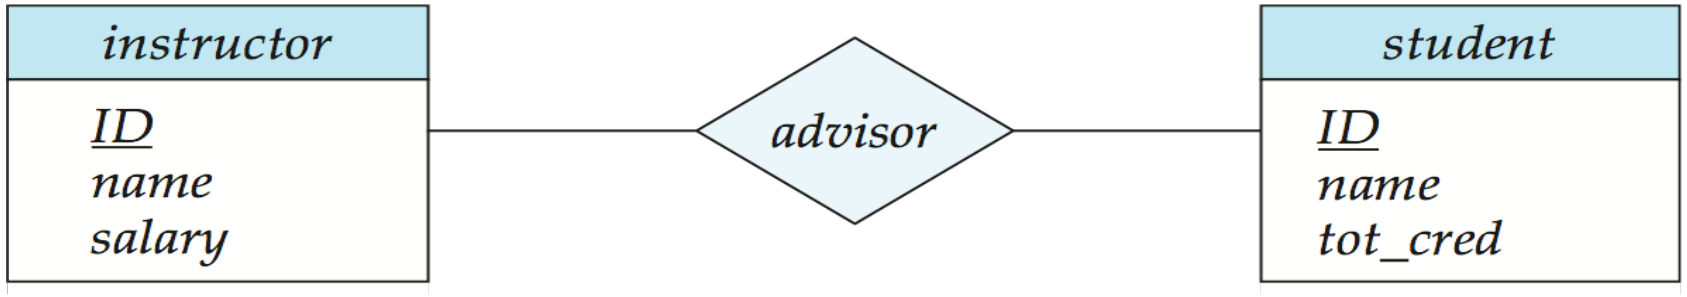

Participation of an Entity Set in a Relationship Set

- **Total participation** (indicated by **double line**): every entity in the entity set participates in at least one relationship in the relationship set
  - E.g., participation of ***section*** in ***sec_course*** is total
    - every ***section*** must have an associated ***course***

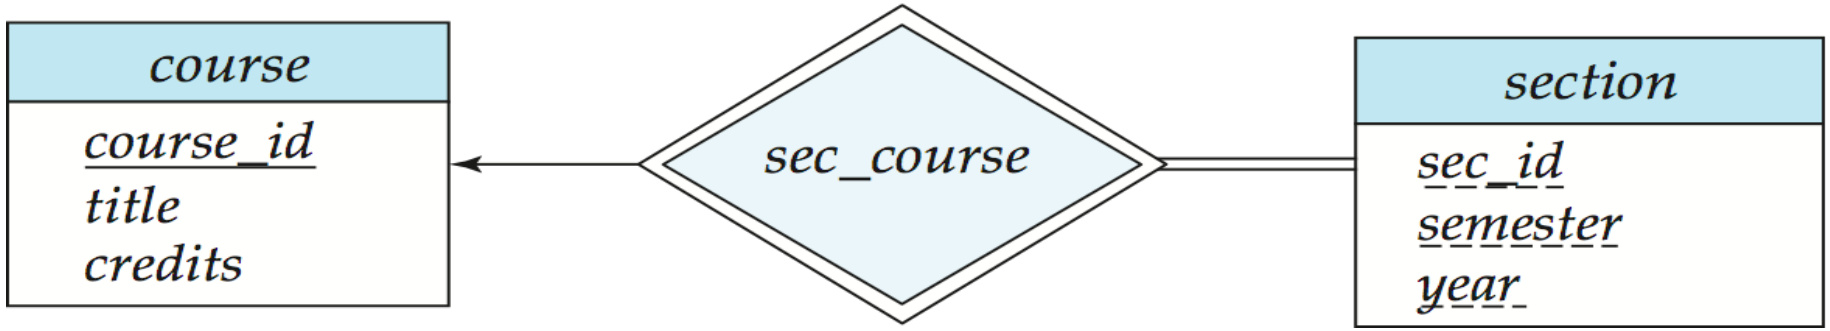

- **Partial participation**: some entities may not participate in any relationship in the relationship set
  - Example: participation of ***instructor*** in ***advisor*** is partial

Alternative Notation for Cardinality Limits

- Cardinality limits can also express participation constraints:
  - The line with cardinality ***1…1*** between *advisor* and *student* means each student must have exactly one advisor.
  - The line with cardinality ***0…**** between *advisor* and *instructor* means an instructor can have zero or more students.
  - The mapping cardinality is **one-to-many** from *instructor* to *student*.
  - The participation of *student* in *advisor* is **total**.

Weak Entity Sets

- An entity set that does not have a primary key is referred to as a **weak entity set**.
- The existence of a weak entity set depends on the existence of a **identifying entity set**
  - It must relate to the identifying entity set via a total, one-to-many relationship set from the identifying to the weak entity set
  - **Identifying relationship** depicted using a **double diamond**
- The **discriminator** (*or partial key)* of a weak entity set is the set of attributes that distinguishes among all the entities of a weak entity set.
- The primary key of a weak entity set is formed by the primary key of the **strong entity set** on which the weak entity set is existence dependent, plus the weak entity set's discriminator.

Weak Entity Sets (Cont.)

- We **underline** the **discriminator** of a weak entity set with a **dashed line**.
- We put the identifying relationship of a weak entity in a double diamond.
- Primary key for *section* – (*course_id, sec_id, semester, year*)

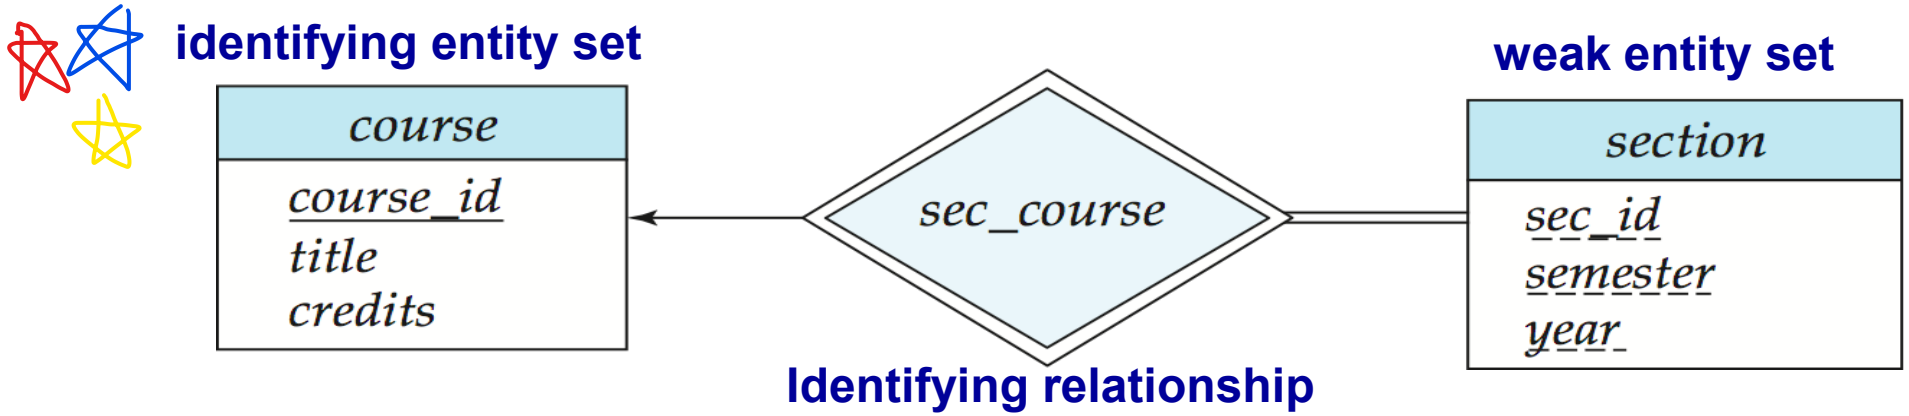

Weak Entity Sets (Cont.)

- Note: the primary key of the **strong entity set** is not explicitly stored with the **weak entity set**, since it is implicit in the **identifying relationship**.

course_id	*sec_id*	*semester*	*year*	*building*	*room_number*	*time_slot_id*
BIO-101	1	Summer	2009	Painter	514	B
BIO-301	1	Summer	2010	Painter	514	A
CS-101	1	Fall	2009	Packard	101	H
CS-101	1	Spring	2010	Packard	101	F
CS-190	1	Spring	2009	Taylor	3128	E
CS-190	2	Spring	2009	Taylor	3128	A
CS-315	1	Spring	2010	Watson	120	D
CS-319	1	Spring	2010	Watson	100	B
CS-319	2	Spring	2010	Taylor	3128	C
CS-347	1	Fall	2009	Taylor	3128	A
EE-181	1	Spring	2009	Taylor	3128	C
FIN-201	1	Spring	2010	Packard	101	B
HIS-351	1	Spring	2010	Painter	514	C
MU-199	1	Spring	2010	Packard	101	D
PHY-101	1	Fall	2009	Watson	100	A

E-R Diagram for University Enterprise

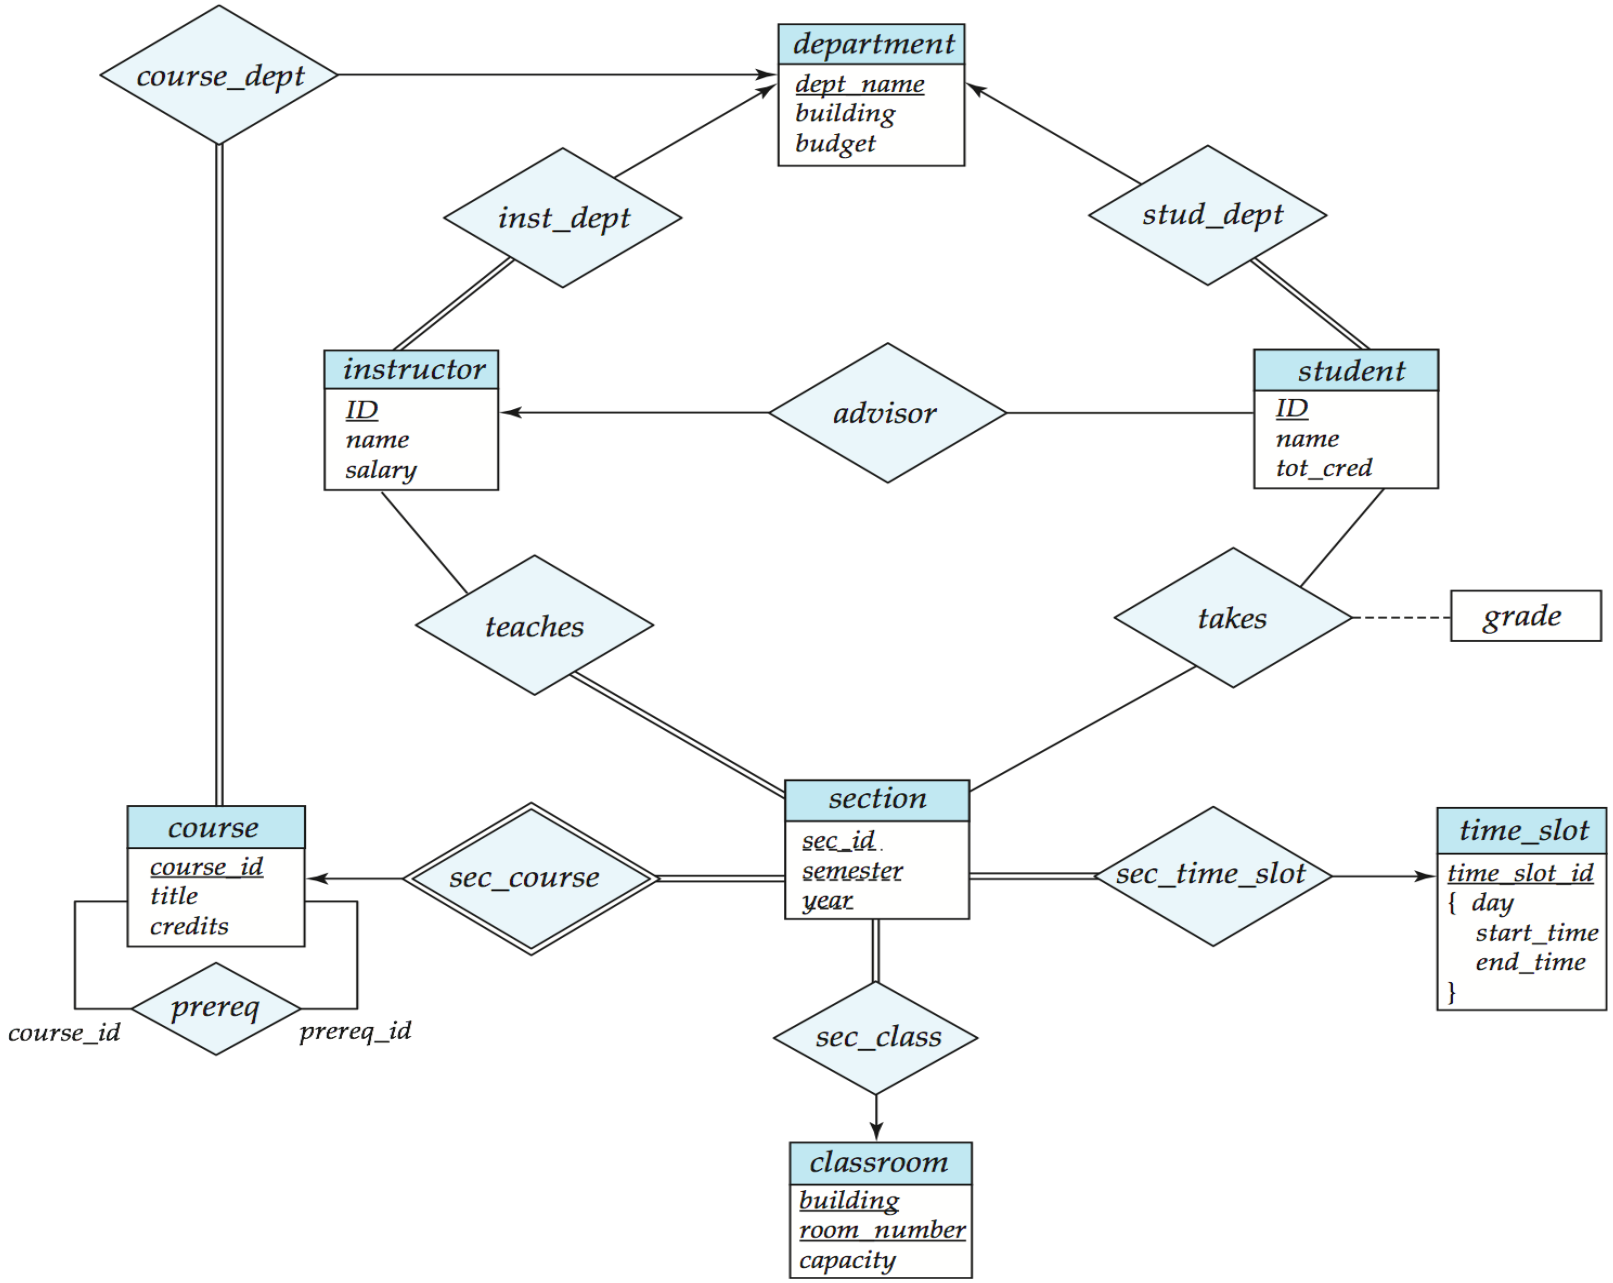

Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as ***relation schemas*** that represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.

Representing Relationship Sets

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*

  ***advisor* = (*instructor_ID*, *student_ID*)**

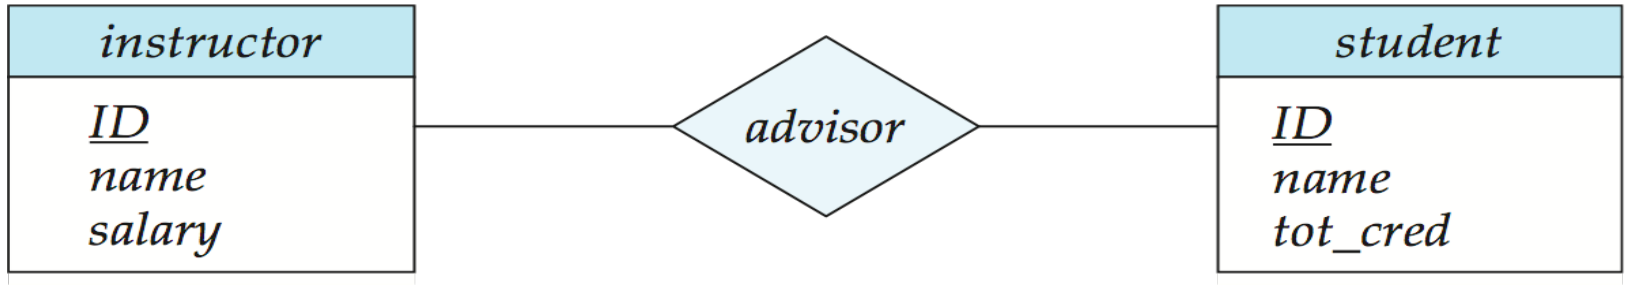

Representing Entity Sets With Simple Attributes

- A strong entity set reduces to a schema with the same attributes
  course = (course_id, title, credits)
- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set ***section = (course_id, sec_id, semester, year)***

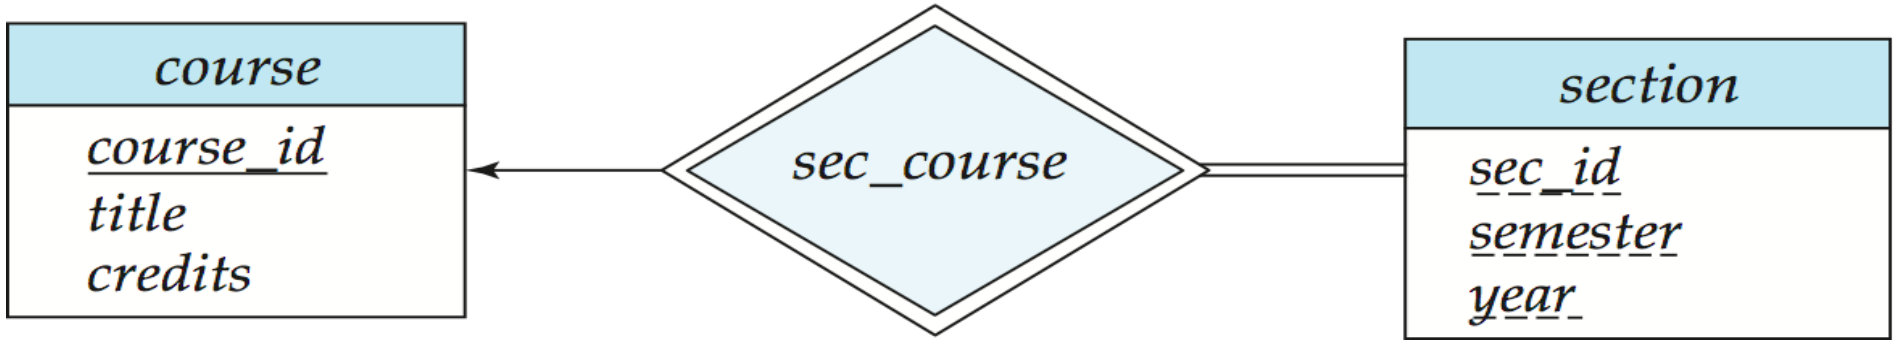

Redundancy of Schemas

- Many-to-one and one-to-many relationship sets that are **total** on the many-side can be represented by adding an extra attribute to the "many" side, containing the primary key of the "one" side
  - Example: Instead of creating a schema for relationship set *inst_dept*, add an attribute *dept_name* to the schema arising from entity set *instructor*

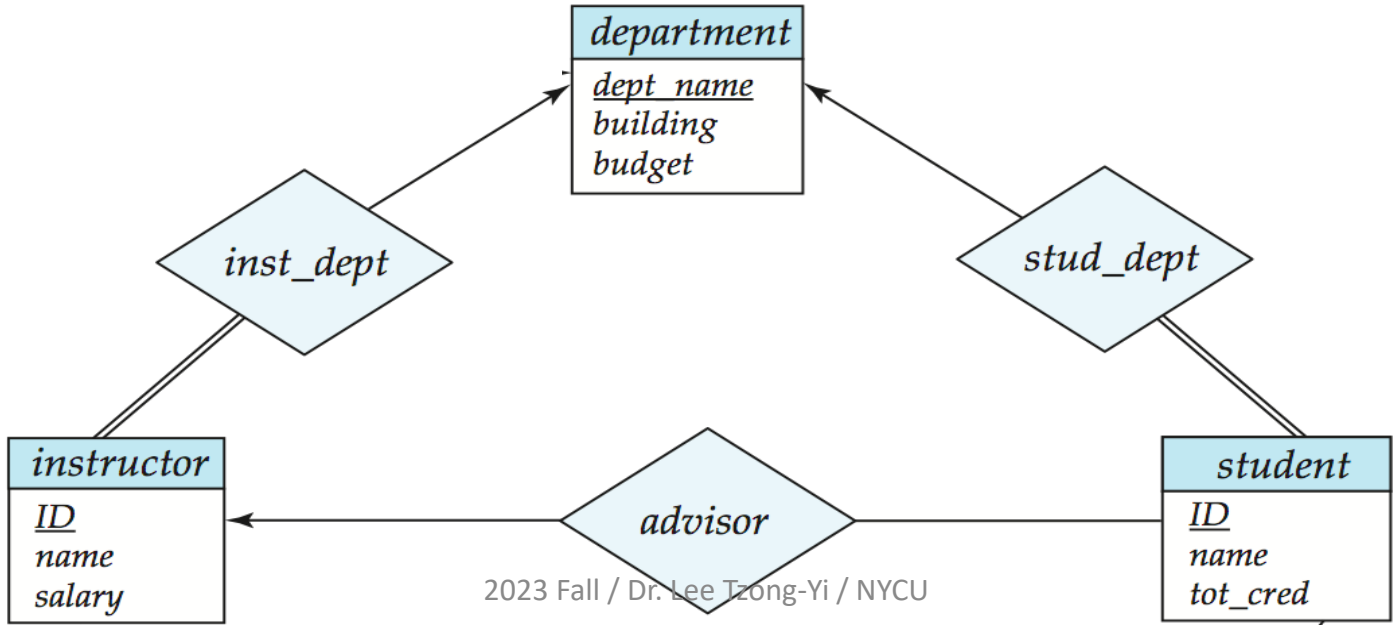

Redundancy of Schemas (Cont.)

- For **one-to-one** relationship sets, either side can be chosen to act as the "many" side
  - That is, extra attribute can be added to either of the tables corresponding to the two entity sets
- If participation is ***partial*** on the "many" side, replacing a schema by an extra attribute in the schema corresponding to the "many" side could result in **null** values
- The schema corresponding to a relationship set linking a weak entity set to its identifying strong entity set is **redundant**.
  - Example: The ***section*** schema already contains the attributes that would appear in the ***sec_course*** schema

Composite and Multivalued Attributes

- Composite attributes are flattened out by creating a separate attribute for each component attribute
  - Example: given entity set ***instructor*** with composite attribute ***name*** containing component attributes ***first_name***, ***middle_initial***, and ***last_name***
  - The schema has three attributes ***first_name***, ***middle_initial***, and ***last_name***
- Ignoring multivalued attributes, extended instructor schema is

  phone number is not here

  - *instructor(ID, first_name, middle_initial, last_name, street_number, street_name, apt_number, city, state, zip_code, date_of_birth)*

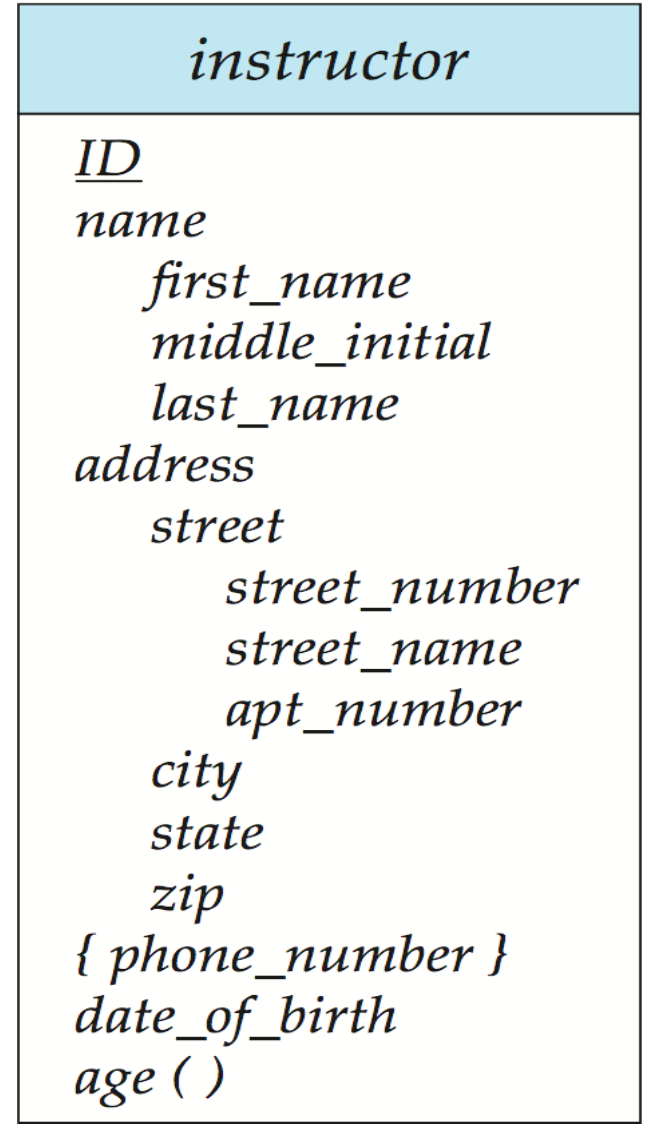

Composite and Multivalued Attributes

- A multivalued attribute *M* of an entity *E* is represented by a separate schema *EM*
  - Schema *EM* has attributes corresponding to the primary key of *E* and an attribute corresponding to multivalued attribute *M*
  - Example: Multivalued attribute ***phone_number*** of ***instructor*** is represented by a schema:

    inst_phone= **(** ***ID, phone_number*** **)**

  - Each value of the multivalued attribute maps to a separate tuple of the relation on schema *EM*
    - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:
      (22222, 456-7890) and (22222, 123-4567)

Design Issues

there is no certain answer

- **Use of entity sets vs. attributes**

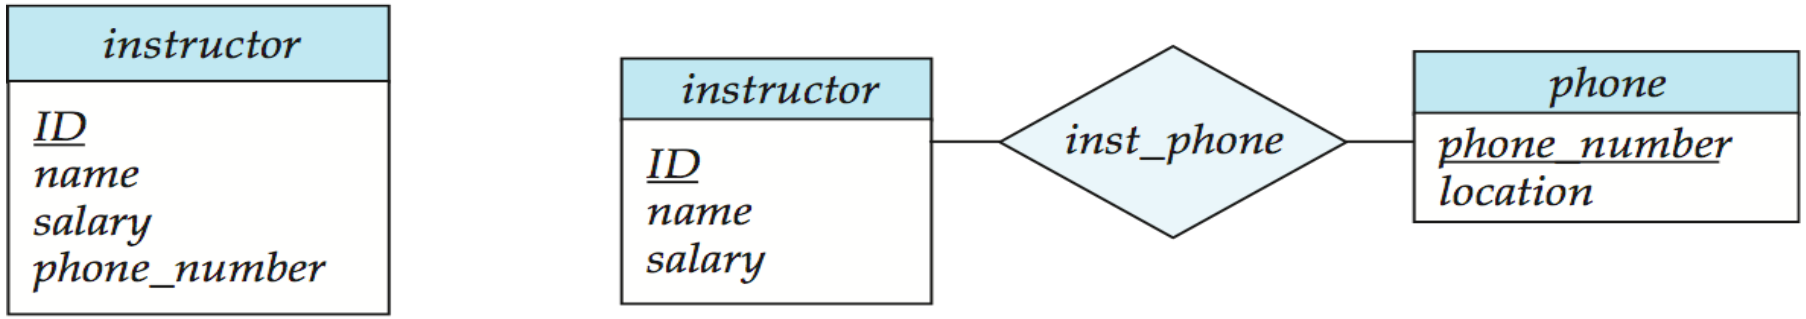

- Use of phone as an entity allows extra information about phone numbers (plus multiple phone numbers)

Design Issues

- **Binary vs. n-ary relationship sets**
  Although it is possible to replace any non-binary (*n*-ary, for $n > 2$) relationship set by a number of distinct binary relationship sets, a *n*-ary relationship set shows more clearly that several entities participate in a single relationship.

- **Placement of relationship attributes**

  e.g., attribute ***date*** as attribute of ***advisor*** or as attribute of ***student***

Converting Non-Binary Relationships to Binary Form

- In general, any non-binary relationship can be represented using binary relationships by creating an artificial entity set.
  - Replace *R* between entity sets A, B and C by an entity set *E*, and three relationship sets:

    1. R_A, relating *E* and *A* 2. R_B, relating *E* and *B* 3. R_C, relating *E* and *C*
  - Create a special identifying attribute for *E*
  - Add any attributes of *R* to *E*
  - For each relationship (a_i , b_i , c_i) in *R,* create

    1. a new entity e_i in the entity set *E* 2. add (e_i , a_i) to R_A

    3. add (e_i , b_i) to R_B 4. add (e_i , c_i) to R_C

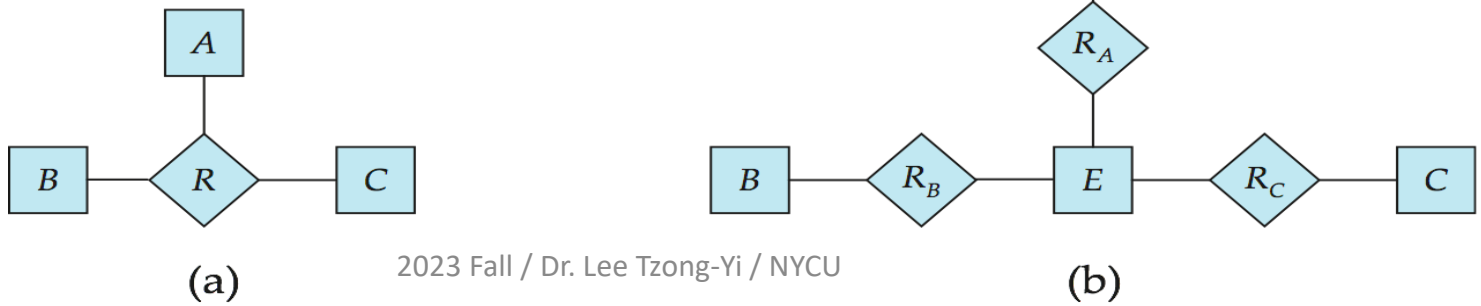

Extended E-R Features: Specialization

- **Top-down design process**: we designate subgroupings within an entity set that are distinctive from other entities in the set.
- These subgroupings become lower-level entity sets that have attributes or participate in relationships that do not apply to the higher-level entity set.
- Depicted by a *triangle* component labeled **ISA** (E.g., *instructor* "is a" *person*).
- **Attribute inheritance** – a lower-level entity set **inherits** all the attributes and relationship participation of the higher-level entity set to which it is linked.

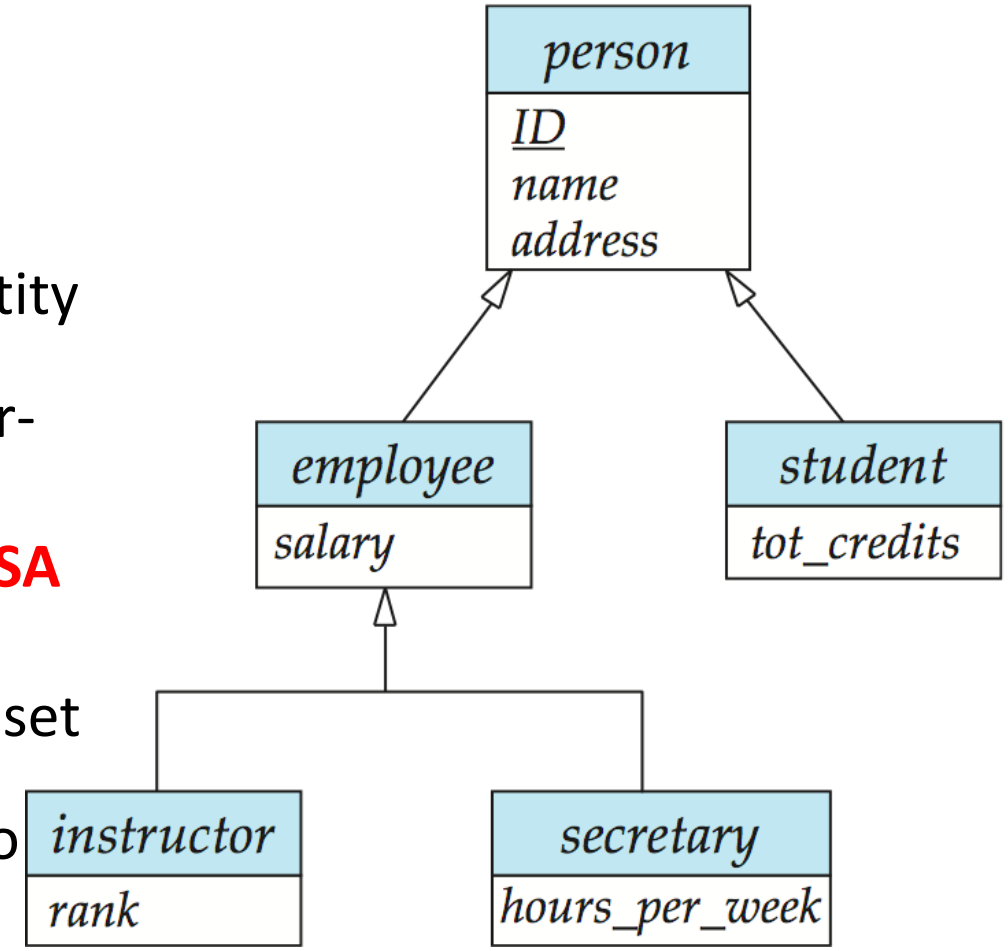

Extended ER Features: Generalization

- A **bottom-up design process** – combine a number of entity sets that share the same features into a higher-level entity set.
- Specialization and generalization are simple inversions of each other; they are represented in an E-R diagram in the same way.
- The terms **specialization** and **generalization** are used interchangeably.

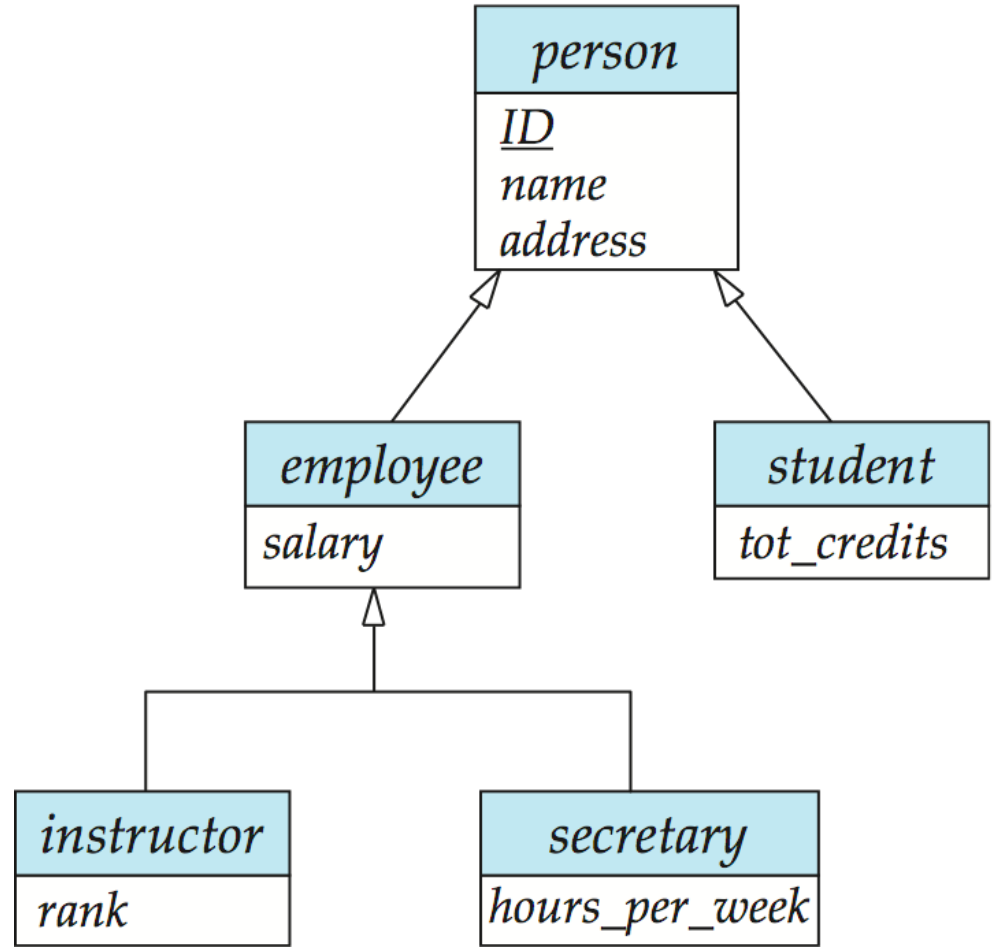

Specialization and Generalization

- Can have **multiple specializations** of an entity set based on different features.
- E.g., *permanent_employee* vs. *temporary_employee*, in addition to *instructor* vs. *secretary*
- Each particular employee would be
  - a member of one of *permanent_employee* or *temporary_employee*,
  - and also a member of one of *instructor*, *secretary*
- The ISA relationship also referred to as **superclass - subclass** relationship

Representing Specialization via Schemas

- Method 1:
  - Form a schema for the higher-level entity
  - Form a schema for each lower-level entity set, include **primary key** of higher-level entity set and local attributes

schema	attributes
person	*ID, name, street, city*
student	*ID, tot_cred*
employee	*ID, salary*

  - **Drawback**: getting information about, an *employee* requires accessing two relations, the one corresponding to the low-level schema and the one corresponding to the high-level schema

Representing Specialization as Schemas (Cont.)

- Method 2:
  - Form a schema for each entity set with all local and inherited attributes

schema	attributes
person	*ID, name, street, city*
student	*ID, name, street, city, tot_cred*
employee	*ID, name, street, city, salary*

  - If specialization is total, the schema for the generalized entity set (*person*) not required to store information
    - Can be defined as a "view" relation containing union of specialization relations
    - But explicit schema may still be needed for foreign key constraints
  - **Drawback**: *name, street* and *city* may be stored redundantly for people who are both students and employees

E-R Design Issues

- The use of **an attribute** or **entity set** to represent an object.
- Whether a real-world concept is best expressed by an **entity set** or a **relationship set**.
- The use of a ternary relationship versus a pair of binary relationships.
- The use of a strong or weak entity set.
- The use of specialization/generalization – contributes to **modularity** in the design.

Summary of Symbols Used in E-R Notation

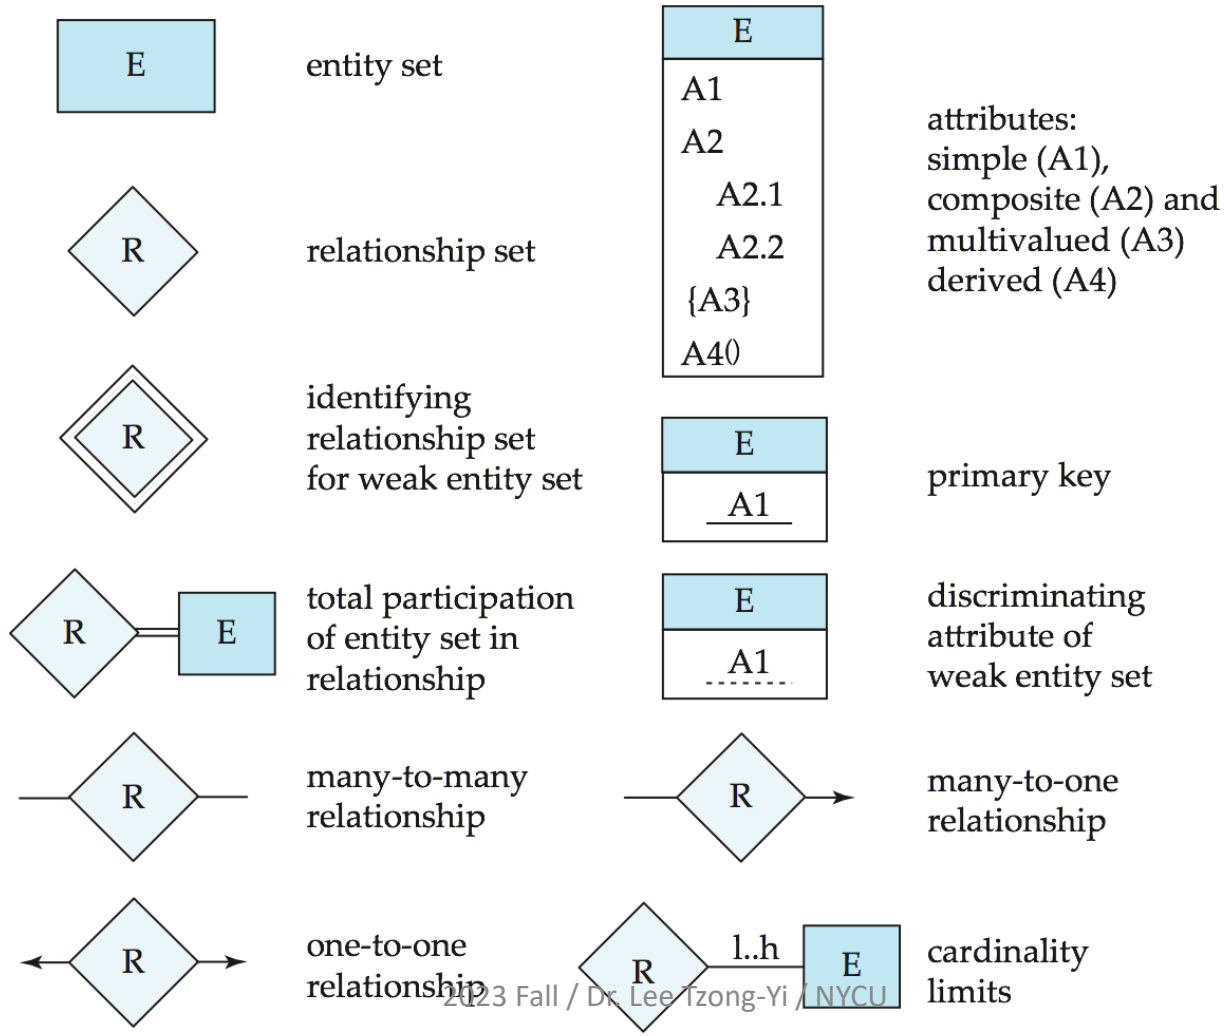

Thank you and any questions?